

→ 03/10/2024

# Optimization

---

DAT-300



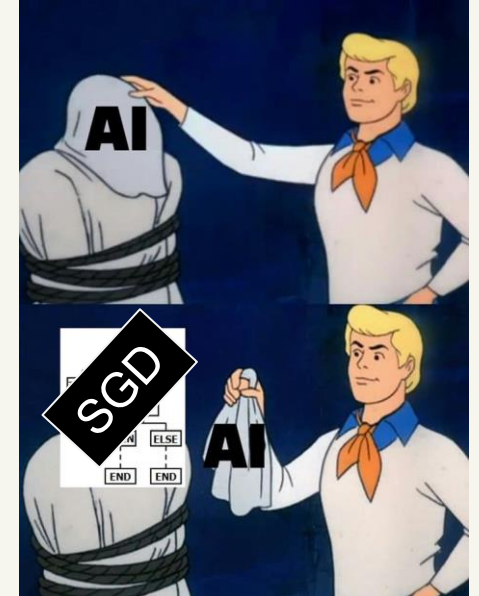
# Agenda

- Motivation
- Key Concepts
- Optimization Algorithms
- Challenges
- Hands on-application

# Motivation ( $\theta$ refers to $w$ in the jupyter notebook)

- Optimization: finding the minimum/maximum value of a function
  - In ML: by adjusting model parameters to minimize an objective function (cost function, loss function) approximated by the data.
  - Approximation: local minimum that perform “well enough” → better generalization?
- “... all of deep learning is powered by one very important algorithm: stochastic gradient descent (1869/1952)” - Ian Goodfellow

$$\theta^* \in \underset{\theta \in \Theta}{\operatorname{argmin}} \mathcal{L}(\theta)$$



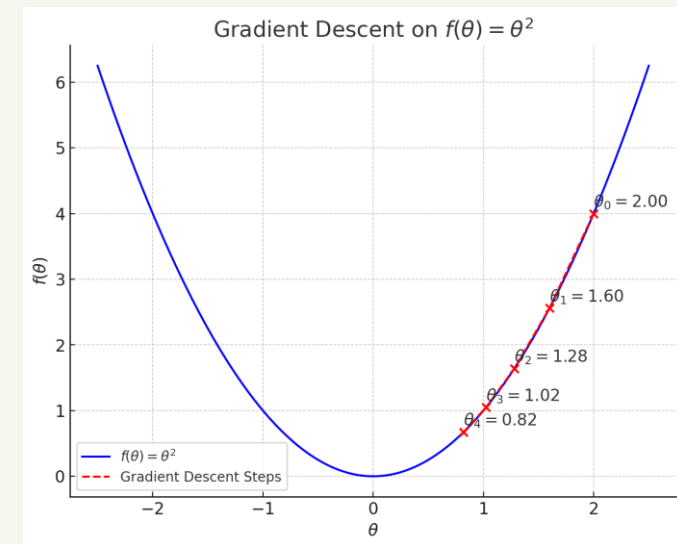
# Key Concepts

- Backpropagation  $\neq$  optimization
  - Computational graph  $\rightarrow$  backward pass
  - Gradients; partially derivatives (rate of change)
- Object (cost, loss) function
- Convex vs Non-convex
- Loss Landscape
  - Min, max, saddle point
- Smooth vs nonsmooth
  - Continuous differentiable
  - [Lipschitz continuity](#)
- Jacobian matrix
  - Generalization to a vector valued function
- Hessian matrix
  - Second order derivative matrix
  - Explains the curvature
  - Newton Method
    - $\rightarrow n^2$  - elements



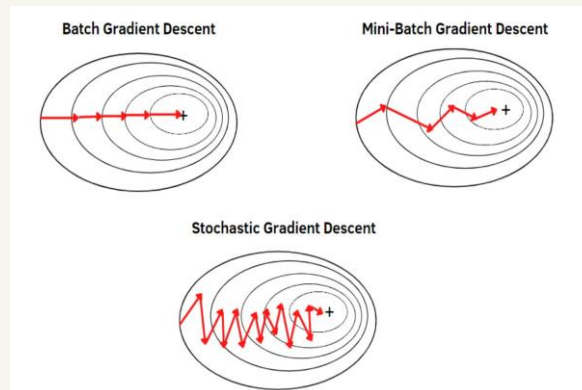
# Gradient Descent - ( $\theta$ refers to $w$ in the jupyter notebook)

- An iterative optimization algorithm to find the/a minimum of a function
- Minimize discrepancy between prediction and true label/value
- Move in the the direction of the steepest decrease (opposite of the Gradient)
- I.e. using MSE
  - $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N [(h_{\theta}(x^i) - y^{(i)})^2]$
  - e.g. linear regression  $\rightarrow h_{\theta}(x^{(i)}) = wx^{(i)} + b$
  - *Repet until converge (treshold is reach)*  
$$\theta \leftarrow \theta_{old} - \eta \nabla \mathcal{L}(\theta)$$
  - Learning rate:  $\eta$



# Stochastic Gradient Descent (SDG)

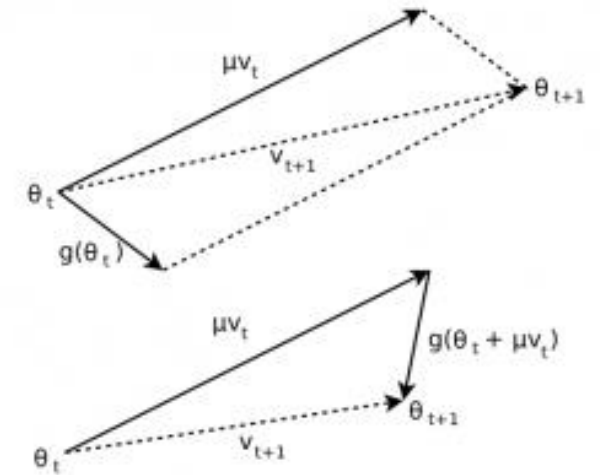
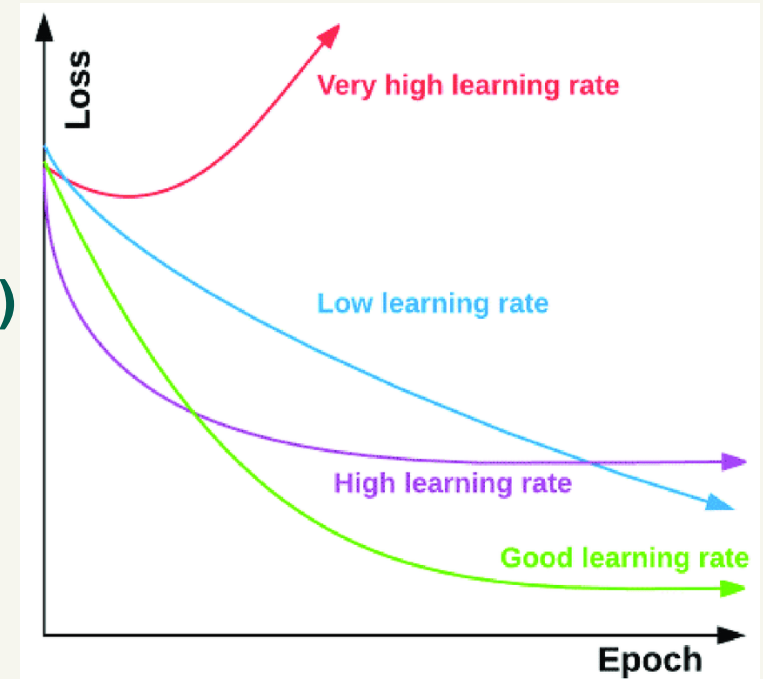
- Problem with GD?
  - $\frac{(3 \times 224 \times 224) \times 4 \times 10000 \text{ samples}}{1024 \times 3} \approx 5.5GB$
- Batch size  $m \ll N$ 
  - $m = 1$  [Stochastic GD]
  - $m = 2^x$  [Mini-batch GD (aka SGD in ML)]
  - $m = N$  [(Batch) GD]
- Error in computing mean  $\rightarrow \frac{\sigma}{\sqrt{m}}$ 
  - **Diminishing returns** - the more samples you use, the smaller the error becomes, but the improvement slows down as  $m$  increases



	Small batch size	Large batch size
Gradient estimate	Noisy (high variance)	Smooth (low variance)
Convergence speed	Fast (unstable)	Slow (stable)
Memory consumption	Low	High
<b>Generalization</b>	Overfit (sample importance)	Better
Learning rate	Small	Higher
Updates	Frequent	Fewer
Common batch sizes	32, 64, 128, (256, 512)	> 2048

# Momentum - ( $\theta$ refers to $w$ in the jupyter notebook)

- Learning rate  $\Rightarrow \theta \leftarrow \theta_{old} - \eta \nabla \mathcal{L}(\theta)$
- Problem
  - Oscillation (when steep), slow convergence (when flat), stuck in local minimum
- Momentum  $\Rightarrow mass \times velocity$
- Update rule:
  - $v_t = \gamma v_{t-1} + \eta \nabla \mathcal{L}(\theta_{t-1})$
  - $\theta \leftarrow \theta_{t-1} - v_t$ 
    - $\rightarrow$  Moment coefficient  $\rightarrow \gamma [0.9, 0.99]$
- Moving average of past gradients  $\rightarrow$  additional hyperparameter
- Nesterov Momentum (“look ahead”)  $\rightarrow v_t = \gamma v_{t-1} + \eta \nabla \mathcal{L}(\theta_{t-1}, v_t)$



# ML Optimization Algorithm - ( $\theta$ refers to $w$ in the jupyter notebook)

“...researchers have long realized that the **learning rate** ... the hyperparameters that is the **most difficult to set** ... a **significant impact** on model performance.” Ian Goodfellow

Solution  $\Rightarrow$  SGD w/ adaptive Learning rate

- AdaGrad  $\rightarrow$  **Adaptive Gradient**
- RMSProp  $\rightarrow$  **Root Mean Square Propagation**
- Adam  $\rightarrow$  **Adaptive Moment (Estimation)**

$$\text{RMSProp: } G_t = \sqrt{\frac{1}{t} \sum_{\tau=1}^t g_{\tau} g_{\tau}^T}$$

Adam: RMSProp + Momentum

AdaGrad:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\epsilon I + \text{diag}(G_t)}} \cdot g_t, \quad | \quad G_t = \sum_{\tau=1}^t g_{\tau} g_{\tau}^T.$$

$$\begin{bmatrix} \theta_{t+1}^{(1)} \\ \theta_{t+1}^{(2)} \\ \vdots \\ \theta_{t+1}^{(m)} \end{bmatrix} = \begin{bmatrix} \theta_t^{(1)} \\ \theta_t^{(2)} \\ \vdots \\ \theta_t^{(m)} \end{bmatrix} - \eta \left( \begin{bmatrix} \epsilon & 0 & \dots & 0 \\ 0 & \epsilon & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \epsilon \end{bmatrix} + \begin{bmatrix} G_t^{(1,1)} & 0 & \dots & 0 \\ 0 & G_t^{(2,2)} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & G_t^{(m,m)} \end{bmatrix} \right)^{-1/2} \begin{bmatrix} g_t^{(1)} \\ g_t^{(2)} \\ \vdots \\ g_t^{(m)} \end{bmatrix} \quad (4)$$

$$\begin{bmatrix} \theta_{t+1}^{(1)} \\ \theta_{t+1}^{(2)} \\ \vdots \\ \theta_{t+1}^{(m)} \end{bmatrix} = \begin{bmatrix} \theta_t^{(1)} \\ \theta_t^{(2)} \\ \vdots \\ \theta_t^{(m)} \end{bmatrix} - \begin{bmatrix} \frac{\eta}{\sqrt{\epsilon + G_t^{(1,1)}}} & 0 & \dots & 0 \\ 0 & \frac{\eta}{\sqrt{\epsilon + G_t^{(2,2)}}} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \frac{\eta}{\sqrt{\epsilon + G_t^{(m,m)}}} \end{bmatrix} \begin{bmatrix} g_t^{(1)} \\ g_t^{(2)} \\ \vdots \\ g_t^{(m)} \end{bmatrix} \quad (5)$$

$$\begin{bmatrix} \theta_{t+1}^{(1)} \\ \theta_{t+1}^{(2)} \\ \vdots \\ \theta_{t+1}^{(m)} \end{bmatrix} = \begin{bmatrix} \theta_t^{(1)} \\ \theta_t^{(2)} \\ \vdots \\ \theta_t^{(m)} \end{bmatrix} - \begin{bmatrix} \frac{\eta}{\sqrt{\epsilon + G_t^{(1,1)}}} g_t^{(1)} \\ \frac{\eta}{\sqrt{\epsilon + G_t^{(2,2)}}} g_t^{(2)} \\ \vdots \\ \frac{\eta}{\sqrt{\epsilon + G_t^{(m,m)}}} g_t^{(m)} \end{bmatrix}$$

SGD:

$$\begin{bmatrix} \theta_{t+1}^{(1)} \\ \theta_{t+1}^{(2)} \\ \vdots \\ \theta_{t+1}^{(m)} \end{bmatrix} = \begin{bmatrix} \theta_t^{(1)} \\ \theta_t^{(2)} \\ \vdots \\ \theta_t^{(m)} \end{bmatrix} - \begin{bmatrix} \eta g_t^{(1)} \\ \eta g_t^{(2)} \\ \vdots \\ \eta g_t^{(m)} \end{bmatrix}$$



# tf.keras.optimizers

```
1 import tensorflow as tf
2
3 # Stochastic gradient descent, with or without momentum
4 op1=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.0,
5                               nesterov=False, name='SGD', **kwargs)
6
7 op2=tf.keras.optimizers.RMSprop(learning_rate=0.001, rho=0.9,
8                                   momentum=0.0, epsilon=1e-07,
9                                   centered=False, name='RMSprop', **kwargs)
10
11 op3=tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9,
12                               beta_2=0.999, epsilon=1e-07,
13                               amsgrad=False, name='Adam', **kwargs)
```

```
8 # Train step for a single batch
9 def train_step(image, y):
10
11     # get predictions loss and update model parameters
12     with tf.GradientTape() as tape:
13         y_pred = model(image, training=True)
14         loss = loss_fn(y, y_pred)
15         grads = tape.gradient(loss, model.trainable_variables)
16         optimizer_fn.apply_gradients(zip(grads, model.
17                                         trainable_variables))
```

```
1 model.compile(
2     optimizer=optimizer_fn, # Optimizer
3     loss=loss_fn, # Loss function to minimize
4     metrics=[metric_fn], # List of metrics to monitor
5 )
6
7 history = model.fit(
8     train_dataset,
9     batch_size=64,
10    epochs=2,
11    validation_data=val_dataset,
12 )
```

# Strategies to improve optimization

## In the algorithm

1. Parameter initialization
  - `keras.kernel_initializer()`
2. Learning rate scheduling/decay
  - `keras.callbacks.LearningRateScheduler()`
3. Early stopping

## In constructing the model

1. Activation function
  - Sigmoid, ReLU, tanh ...  $\Rightarrow$  vanishing and exploding gradient
2. Batch Normalization
  - Stable gradient
  - Deeper network  $\Rightarrow$  sensitive to input distribution

