

# Sentiment Analysis

## TF-IDF

**Term Frequency (TF)** — How many times this word appears in a document (sentence)  
**Inverse Document Frequency (IDF)** — the natural logarithm of the total number of documents divided by (1 + the total number of documents that contain this certain word)

**Example:**

```
documents = [  
    "apple apple",  
    "apple orange"  
]
```

$$\text{idf}(t, d) = \log \frac{n_d}{1 + \text{df}(d, t)},$$

**Term Frequency (TF)**

	apple	orange
0	2	0
1	1	1

**IDF**

$$\log \frac{2}{1+2} = \log = \log \frac{2}{3}$$

$$\log \frac{2}{1+1} = \log = \log \frac{2}{2}$$

# Sentiment Analysis

## TF-IDF

**Term Frequency (TF)** — How many times this word appears in a document (sentence)  
**Inverse Document Frequency (IDF)** — the natural logarithm of the total number of documents divided by (1 + the total number of documents that contain this certain word)

**Example:**

```
documents = [  
    "apple apple",  
    "apple orange"  
]
```

TF-IDF: The higher the TF-IDF score of a word, the more relevant this word is to the document.  
Orange is more relevant to the document than *apple* here.

### Term Frequency (TF)

	apple	orange
0	2	0
1	1	1

### IDF

$$\log \frac{2}{1+2} = \log = \log \frac{2}{3}$$

$$\log \frac{2}{1+1} = \log = \log \frac{2}{2}$$

# Sentiment Analysis

## TF-IDF

**Term Frequency (TF)** — How many times this word appears in a document (sentence)  
**Inverse Document Frequency (IDF)** — the natural logarithm of the total number of documents divided by (1 + the total number of documents that contain this certain word)

**Example:**

```
documents = [  
    "apple apple",  
    "apple orange"  
]
```

TF: the more a word appears, the more relevant the word is to the document

**Term Frequency (TF)**

	apple	orange
0	2	0
1	1	1

**IDF**

$$\log \frac{2}{1+2} = \log = \log \frac{2}{3}$$

$$\log \frac{2}{1+1} = \log = \log \frac{2}{2}$$

# Sentiment Analysis

## TF-IDF

**Term Frequency (TF)** — How many times this word appears in a document (sentence)  
**Inverse Document Frequency (IDF)** — the natural logarithm of the total number of documents divided by (1 + the total number of documents that contain this certain word)

**Example:**

```
documents = [  
    "apple apple",  
    "apple orange"  
]
```

The IDF measures how rare or common a word is in the collection of documents.

- If a word appears in many documents, its IDF score will be lower, as this word is *less unique to this certain document*.
- If a word appears in fewer documents, its IDF score will be higher, as this word is *more unique to this certain document*.

**Term Frequency (TF)**

	apple	orange
0	2	0
1	1	1

**IDF**

$$\log \frac{2}{1+2} = \log = \log \frac{2}{3}$$

$$\log \frac{2}{1+1} = \log = \log \frac{2}{2}$$

# Word2Vec

- Word2Vec is a group of related models that transform words into vectors.
- Captures semantic meaning based on the context of words.

# Word2Vec

- Word2Vec is a group of related models that transform words into vectors.
- Captures semantic meaning based on the context of words.


## *Why?*

- Traditional methods represent words as discrete symbols.
- Need for capturing semantic meaning in a continuous space.



One-hot word vectors:

- Sparse
- High-dimensional
- Hard-coded




Word embeddings:


- Dense
- Lower-dimensional
- Learned from data

# Word2Vec Models

- Continuous Bag of Words (CBOW)
- Skip-Gram



**CBOW**



**Skip-gram**

# Continuous Bag of Words (CBOW)

- **Goal:** Given a context, predict the target word.

## *How?*

- It takes the context (surrounding words) as input and tries to predict the word in the middle. The idea is that the semantic meaning of a word can be inferred by the words surrounding it.
- For **CBOW**, the one-hot encoded context vectors are averaged or summed, passed through the hidden layer, and then the output layer to predict the target word.

## Example

Sentence: "The cat sat on the mat."

- For the word "sat", if we consider a window size of 2, the context words are ["The", "cat", "on", "the"].
- The target word is "sat".
- CBOW tries to predict "sat" given ["The", "cat", "on", "the"].



# Skip-Gram

- **Goal:** Given a word, predict the context.

## *How?*

- It takes a word as input and tries to predict the context, meaning the surrounding words. It's believed to work better with a small amount of data and with rare words.
- For **Skip-Gram**, the one-hot encoded word vector is passed through the hidden layer and then the output layer to predict context words.

## Example

Sentence: "The cat sat on the mat."

- For the word "sat", with a window size of 2, the target words are ["The", "cat", "on", "the"].
- Skip-Gram tries to predict each of these context words given the input word "sat".

# Differences

- **Data Efficiency:** Skip-Gram usually works well with less data and captures more detailed word relationships.
- **Performance:** CBOW is faster since it averages over the context words, but Skip-Gram generally has superior quality, especially with infrequent words.

